

Research Meeting Follow-up Assistant

Persistent meeting memory for research teams

- Built for PhD advisors and their researchers
- Tracks commitments and reminds owners between meetings
- Extends CCB's transcripts with cited history and scheduled follow-up

Engineering question *Can an assistant reliably maintain commitments across a sequence of meetings while reducing manual reconciliation and avoiding unsupported updates?*

Context Gets Lost Between Meetings

- 1 Weekly research meetings generate commitments, decisions, deadlines, and professor suggestions that must be tracked across weeks, and reconciled against everything decided before.
- 2 A meeting summary captures what happened in one meeting. Connecting it back to earlier commitments, deadlines, and open suggestions is left entirely to the user.
- 3 The result: duplicated work, forgotten commitments, unclear ownership, and decisions whose original reasoning becomes hard to recover.

ONE TASK ACROSS THREE MEETINGS (Example):

- Meeting 4** Will commits to rerun the baseline by Friday. Before the deadline, the system sends him a scheduled reminder.
- Meeting 5** “The rerun finished early.” Evidence updates the existing task to done and cancels pending reminders.
- Meeting 6** The advisor's papers remain unsent. The task stays open, appears in the next brief, and receives follow-up on the configured schedule.

Professor suggestions stay separate from commitments until a researcher explicitly accepts them.

Our Solution: How It Works

1 Detect

Watch a configured folder; gate on consent; associate the artifact with a project and calendar event.

3 Extract

Pull decisions, owned commitments, and professor suggestions, each with source evidence.

5 Validate

Apply evidenced changes automatically, or route ambiguous matches to human review.

2 Transcribe & confirm

Import time stamped transcript; confirm or correct speaker mapping.

4 Reconcile

Match new statements against persistent project state Automatically send follow-ups.

6 Answer

Cited historical Q&A over the decision ledger, with transcript timestamps and audio references.

Tech stack	Python service · SQLite · one LLM behind a schema-validated adapter · Whisper / pyannote.audio if the sponsor pipeline can't be reused.
Focus areas	Prompting + RAG for cited Q&A. CLI and local service. Offline tests and follow-up delivery checks.
Deferred	Additional delivery channels, native reminder sync, advanced diarization, voice enrollment, backward planning.

Evidence, Safeguards, and Non-Goals

EVIDENCE & REVIEW

- Every participant consents before recording or processing; a non-recorded option is always available.
- Speaker mappings are user-confirmed; every task or decision change cites its source segment and timestamp.
- Review ambiguous matches. Send follow-ups only to user-configured recipients on an enabled schedule.
- Silence in a later meeting never marks a task done or dropped; partial progress is not completion.

PRIVACY & NON-GOALS

- Data stays local by default; deletion removes derived artifacts (transcripts, embeddings, indexes), not just audio.
- Retrieval and reconciliation are scoped per project - one project's data never appears in another's results.
- No productivity or performance scores, no ranking of students or researchers, no monitoring without consent.
- Sponsor code, datasets, and models are used only after confirming license and redistribution terms.

Use consenting team recordings and staged scenarios; confirm source licenses and provider terms before use.

Evaluation

Does persistent memory beat independent meeting extraction?

Pilot: 6 meetings in six sequences
Held-out human annotations

≥95% update precision
Zero observed unsupported completions

Score tasks, summaries, and cited answers.
Test reminder timing, cancellation, and duplicates.

Provisional targets. No measured results yet.

Delivery plan and feedback

PROPOSED OWNERS*

Will: extraction and reconciliation

Bryan: runtime, CLI, and integrations

Guadalupe: data and evaluation

Week 5: transcript pipeline Week 8: reconciliation

Weeks 11–12: held-out evaluation

Proposed budget: \$75, using existing laptops

Is the three-meeting scope sufficient?

Can the sponsor's transcript pipeline remain upstream?